

# Collaborative Filtering Recommendation System

## Overview

This project focuses on building a collaborative filtering recommendation system using PySpark and Google Cloud Dataproc. Collaborative filtering predicts the interests of a user by collecting preferences from many users.

## Problem Statement

The goal is to provide accurate movie recommendations by predicting user ratings for movies they haven't watched yet, based on historical user-movie interaction data.

## Scope

The project involves:

- Data preparation
- Model training using Alternating Least Squares (ALS) in PySpark
- Evaluation of the model's performance
- Deployment on Google Cloud Dataproc

## Table of Contents

1. Setup the Environment
2. Download Programs and Related Documentation
3. Process of Program Execution
4. Screenshot of Execution Results
5. Conclusion
6. Appendix
7. References

## Setup the Environment

1. **Create a Google Cloud Storage (GCS) Bucket:**

- Create a bucket in GCS to store your scripts and data.

```
gsutil mb gs://movie_recommendation_with_mllib_collaborative_filter
```

- 2.

3. **Upload Data and Scripts to GCS:**

- Upload the movies.csv, ratings.csv, and your PySpark script (Recommendation\_Engine\_MovieLens.py) to your GCS bucket.

```
gsutil cp movies.csv gs://movie_recommendation_with_mllib_collaborative_filter/  
gsutil cp ratings.csv gs://movie_recommendation_with_mllib_collaborative_filter/  
gsutil cp Recommendation_Engine_MovieLens.py  
gs://movie_recommendation_with_mllib_collaborative_filter/
```

- 4.

## Download Programs and Related Documentation

Clone the repository:

```
git clone https://github.com/ASD-Are/Big_Data
```

- 
- Download the MovieLens dataset: [MovieLens 100k Dataset](#)

## Process of Program Execution

### Data Preparation

**Create the u.data File:**

```
vim u.data
```

1. Populate u.data with your data in the format (UserID, MovieID, rating, Timestamp).

**Transform the Raw Data:**

```
cat u.data | tr -s ' ' | cut -d ' ' -f1-3 | tr ' ','' > u_data_transformed.csv
```

2. This command trims extra spaces, extracts the first three fields (UserID, MovieID, rating), and replaces spaces with commas. The transformed data is saved in u\_data\_transformed.csv.

**Upload the Transformed Data:**

```
gsutil cp u_data_transformed.csv gs://movie_recommendation_with_mllib_collaborative_filter/
```

- 3.

## Create PySpark Script

Create a PySpark script `Recommendation_Engine_MovieLens.py` with the following content:

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, explode
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.ml.recommendation import ALS
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
import argparse

# Parse command-line arguments
parser = argparse.ArgumentParser()
parser.add_argument('--input_path_movies', required=True)
parser.add_argument('--input_path_ratings', required=True)
args = parser.parse_args()

# Initialize Spark session
spark = SparkSession.builder.appName('Recommendations').getOrCreate()

# Load data from GCS
movies = spark.read.csv(args.input_path_movies, header=True)
ratings = spark.read.csv(args.input_path_ratings, header=True)

# Preprocess data
ratings = ratings \
    .withColumn('userId', col('userId').cast('integer')) \
    .withColumn('movieId', col('movieId').cast('integer')) \
    .withColumn('rating', col('rating').cast('float')) \
    .drop('timestamp')

# Split data into training and testing sets
(train, test) = ratings.randomSplit([0.8, 0.2], seed=1234)

# Build ALS model
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
          nonnegative=True, implicitPrefs=False, coldStartStrategy="drop")
param_grid = ParamGridBuilder() \
    .addGrid(als.rank, [10, 50, 100, 150]) \
    .addGrid(als.regParam, [.01, .05, .1, .15]) \
    .build()

evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
                               predictionCol="prediction")
cv = CrossValidator(estimator=als, estimatorParamMaps=param_grid,
                    evaluator=evaluator, numFolds=5)
```

```

# Train model
model = cv.fit(train)
best_model = model.bestModel

# Evaluate model
test_predictions = best_model.transform(test)
RMSE = evaluator.evaluate(test_predictions)
print(f"Root-mean-square error = {RMSE}")

# Generate recommendations
nrecommendations = best_model.recommendForAllUsers(10)
nrecommendations = nrecommendations \
    .withColumn("rec_exp", explode("recommendations")) \
    .select('userId', col("rec_exp.movieId"), col("rec_exp.rating"))
nrecommendations.show()

# Join with movie titles for better interpretability
nrecommendations.join(movies, on='movieId').filter('userId = 100').show()
ratings.join(movies, on='movieId').filter('userId = 100').sort('rating',
                                                                ascending=False).limit(10).show()

# Stop Spark session
spark.stop()

```

•

## Create Dataproc Cluster

**Create the cluster with the desired configuration:**

```

gcloud dataproc clusters create spark-cluster \
  --region us-west1 \
  --zone us-west1-a \
  --master-machine-type n1-standard-4 \
  --worker-machine-type n1-standard-4 \
  --num-workers 2

```

1.

## Submit PySpark Job

**Submit the PySpark job to the Dataproc cluster:**

```

gcloud dataproc jobs submit pyspark
gs://movie_recommendation_with_mllib_collaborative_filter/Recommendation_Engine_MovieLe
ns.py --cluster=spark-cluster --region=us-west1 --

```

```
--input_path_movies=gs://movie_recommendation_with_mllib_collaborative_filter/movies.csv  
--input_path_ratings=gs://movie_recommendation_with_mllib_collaborative_filter/ratings.csv
```

1.

## Screenshot of Execution Results

After running the job, the output will display the Root Mean Squared Error (RMSE) of the model.  
Root-mean-square error = 0.48149423210378404

- 

## Conclusion

- Successfully built a collaborative filtering recommendation system using PySpark and Google Cloud Dataproc.
- Data preparation involved transforming and uploading the MovieLens dataset to Google Cloud Storage.
- Model training utilized the ALS algorithm on user-movie interaction data.
- The model's performance was evaluated using Root Mean Squared Error (RMSE), achieving a value of 0.48149423210378404, indicating reasonable prediction accuracy.

## References

- [MovieLens Dataset](#)
- [Collaborative Filtering - RDD-based API](#)