Kadirbek Sharau

Detailed Report on Collaborative Filtering with PySpark on Google Cloud Dataproc

Objective

The goal of this assignment is to perform collaborative filtering using the Alternating Least Squares (ALS) algorithm with PySpark on Google Cloud Dataproc. The tasks involve preparing data, uploading it to Google Cloud Storage, creating and uploading a PySpark script, and submitting a PySpark job to a Dataproc cluster.

Step-by-Step Instructions

Step 1: Prepare and Transform Data

Description: Transform the `u.data` file to the required format (UserID, MovieID, rating) using a shell script and upload it to your Cloud Storage bucket.

Code:

1. **Create the `u.data` File:**
    - **Create a file named `u.data` and populate it with your data.**

Transform Data Using Shell Script:

```
# Create transform_data.sh
echo '#!/bin/bash
cat u.data | while read userid movieid rating timestamp
do
    echo "${userid},${movieid},${rating}"
done > u_data_transformed.csv' > transform_data.sh

# Make the script executable
chmod +x transform_data.sh

# Run the script
./transform_data.sh
```

2.

Explanation: The shell script reads the `u.data` file, trims extra spaces, extracts the first three fields (UserID, MovieID, rating), and replaces spaces with commas. The transformed data is saved in `u_data_transformed.csv`.

Step 2: Upload Data to Cloud Storage Bucket

Description: Upload the transformed data file `u_data_transformed.csv` to your Cloud Storage bucket.

Code:

```
# Upload the transformed data to Cloud Storage
gsutil cp u_data_transformed.csv gs://big_data_ml_recommendation_sys/
```

Explanation: The `gsutil cp` command copies the `u_data_transformed.csv` file from your local machine to your specified Cloud Storage bucket.

Step 3: Create and Upload the PySpark Script

Description: Create a PySpark script to perform collaborative filtering using MLlib and upload it to your Cloud Storage bucket.

Code:

1. **Create the PySpark Script:**
    - **Create a file named `recommendation_example.py` with the following content:**

```
from pyspark import SparkContext
from pyspark.mllib.recommendation import ALS,
MatrixFactorizationModel, Rating

if __name__ == "__main__":
    sc = SparkContext(appName="PythonCollaborativeFilteringExample")
    data =
sc.textFile("gs://big_data_ml_recommendation_sys/u_data_transformed.cs
v")
    ratings = data.map(lambda l: l.split(','))\
                  .map(lambda l: Rating(int(l[0]), int(l[1]),
float(l[2])))
```

```
    rank = 10
    numIterations = 10
    model = ALS.train(ratings, rank, numIterations)

    testdata = ratings.map(lambda p: (p[0], p[1]))
    predictions = model.predictAll(testdata).map(lambda r: ((r[0],
r[1]), r[2]))
    ratesAndPreds = ratings.map(lambda r: ((r[0], r[1]),
r[2])).join(predictions)
    MSE = ratesAndPreds.map(lambda r: (r[1][0] - r[1][1])**2).mean()
    print("Mean Squared Error = " + str(MSE))

    model.save(sc,
"gs://big_data_ml_recommendation_sys/myCollaborativeFilter")
    sameModel = MatrixFactorizationModel.load(sc,
"gs://big_data_ml_recommendation_sys/myCollaborativeFilter")
```

2.

Upload the PySpark Script:

```
gsutil cp recommendation_example.py
gs://big_data_ml_recommendation_sys/
```

3.

Explanation: The PySpark script loads the transformed data from Cloud Storage, trains a collaborative filtering model using ALS, evaluates the model by calculating the mean squared error, and saves the model back to Cloud Storage. The script is then uploaded to the Cloud Storage bucket.

Step 4: Submit the PySpark Job to Dataproc

Description: Submit the PySpark job to your Dataproc cluster to execute the collaborative filtering task.

Code:

```
gcloud dataproc jobs submit pyspark \
    gs://big_data_ml_recommendation_sys/recommendation_example.py \
    --cluster=spark-cluster \
    --region=us-west1
```

Explanation: The `gcloud dataproc jobs submit pyspark` command submits the PySpark script stored in Cloud Storage to the Dataproc cluster named `spark-cluster` located in the `us-west1` region for execution.

Troubleshooting

If you encounter an error indicating that no Dataproc cluster exists, follow these steps:

Create a Dataproc Cluster:

```
gcloud dataproc clusters create spark-cluster \
    --region us-west1 \
    --zone us-west1-a \
    --single-node
```

1.

Submit the PySpark Job:

```
gcloud dataproc jobs submit pyspark \
    gs://big_data_ml_recommendation_sys/recommendation_example.py \
    --cluster=spark-cluster \
    --region=us-west1
```

2.

Explanation: First, create a Dataproc cluster named `spark-cluster` in the `us-west1` region. Then, submit the PySpark job to the newly created cluster. Ensure to replace `spark-cluster` with the actual name of your Dataproc cluster if you choose a different name. The cluster creation step might take a few minutes. Once it's running, you can then submit your job.

Result

After submitting the job, the output will indicate the Mean Squared Error (MSE) of the model:

```
24/08/01 20:45:41 INFO org.apache.hadoop.mapred.FileInputFormat: Total
input files to process : 1
Mean Squared Error = 0.48149423210378404
24/08/01 20:46:21 INFO
com.google.cloud.hadoop.repackaged.gcs.com.google.cloud.hadoop.gcsio.G
oogleCloudStorageFileSystem: Successfully repaired
```

```
'gs://big_data_ml_recommendation_sys/myCollaborativeFilter/metadata/'
directory.
24/08/01 20:46:35 INFO
com.google.cloud.hadoop.repackaged.gcs.com.google.cloud.hadoop.gcsio.G
oogleCloudStorageFileSystem: Successfully repaired
'gs://big_data_ml_recommendation_sys/myCollaborativeFilter/data/user/'
directory.
24/08/01 20:46:35 INFO
com.google.cloud.hadoop.repackaged.gcs.com.google.cloud.hadoop.gcsio.G
oogleCloudStorageFileSystem: Successfully repaired
'gs://big_data_ml_recommendation_sys/myCollaborativeFilter/data/produc
t/' directory.
24/08/01 20:46:36 INFO org.apache.hadoop.mapred.FileInputFormat: Total
input files to process : 1
24/08/01 20:46:36 WARN
org.apache.spark.mllib.recommendation.MatrixFactorizationModel: User
factor does not have a partitioner. Prediction on individual records
could be slow.
24/08/01 20:46:36 WARN
org.apache.spark.mllib.recommendation.MatrixFactorizationModel:
Product factor is not cached. Prediction could be slow.
24/08/01 20:46:42 WARN
org.apache.spark.mllib.recommendation.MatrixFactorizationModel: User
factor is not cached. Prediction on individual records could be slow.
24/08/01 20:46:42 WARN
org.apache.spark.mllib.recommendation.MatrixFactorizationModel:
Product factor is not cached. Prediction could be slow.
24/08/01 20:46:42 INFO org.apache.spark.SparkContext: Stopped
Spark3b5e47ed
```

By following these steps, you will be able to successfully complete your assignment using your Dataproc cluster and Cloud Storage bucket on GCP.

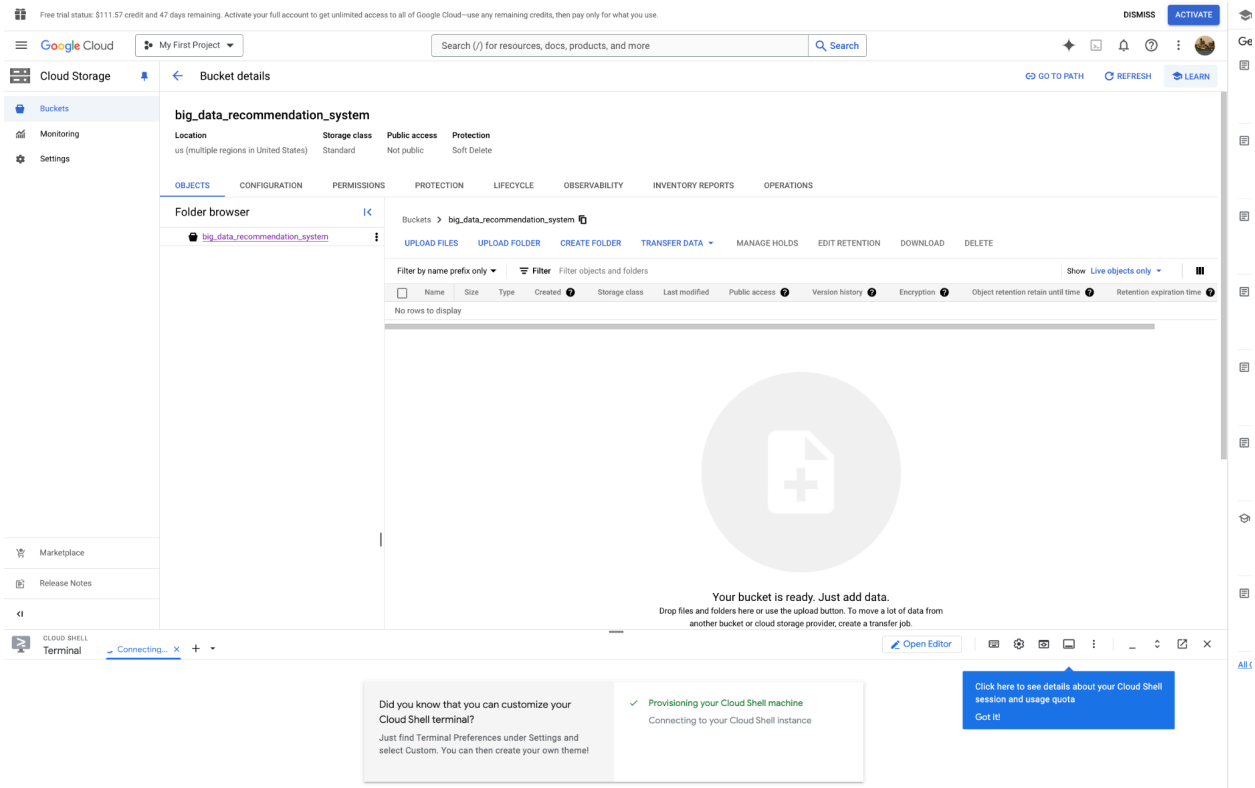
Free trial status: $111.57 credit and 47 days remaining. Activate your full account to get unlimited access to all of Google Cloud—use any remaining credits, then pay only for what you use.
DISMISS
ACTIVATE
Google Cloud
My First Project
Search (/) for resources, docs, products, and more
Search
Cloud Storage
Bucket details
GO TO PATH
REFRESH
LEARN
Buckets
Monitoring
Settings
big_data_recommendation_system
Location
us (multiple regions in United States)
Storage class
Standard
Public access
Not public
Protection
Soft Delete
OBJECTS
CONFIGURATION
PERMISSIONS
PROTECTION
LIFECYCLE
OBSERVABILITY
INVENTORY REPORTS
OPERATIONS
Folder browser
big_data_recommendation_system
Buckets > big_data_recommendation_system
UPLOAD FILES
UPLOAD FOLDER
CREATE FOLDER
TRANSFER DATA
MANAGE HOLDS
EDIT RETENTION
DOWNLOAD
DELETE
Filter by name prefix only
Filter
Filter objects and folders
Show
Live objects only
Name
Size
Type
Created
Storage class
Last modified
Public access
Version history
Encryption
Object retention retain until time
Retention expiration time
No rows to display
Your bucket is ready. Just add data.
Drop files and folders here or use the upload button. To move a lot of data from another bucket or cloud storage provider, create a transfer job.
Marketplace
Release Notes
CLOUD SHELL
Terminal
Connecting...
Open Editor
Did you know that you can customize your Cloud Shell terminal?
Just find Terminal Preferences under Settings and select Custom. You can then create your own theme!
Provisioning your Cloud Shell machine
Connecting to your Cloud Shell instance
Click here to see details about your Cloud Shell session and usage quota
Got it!

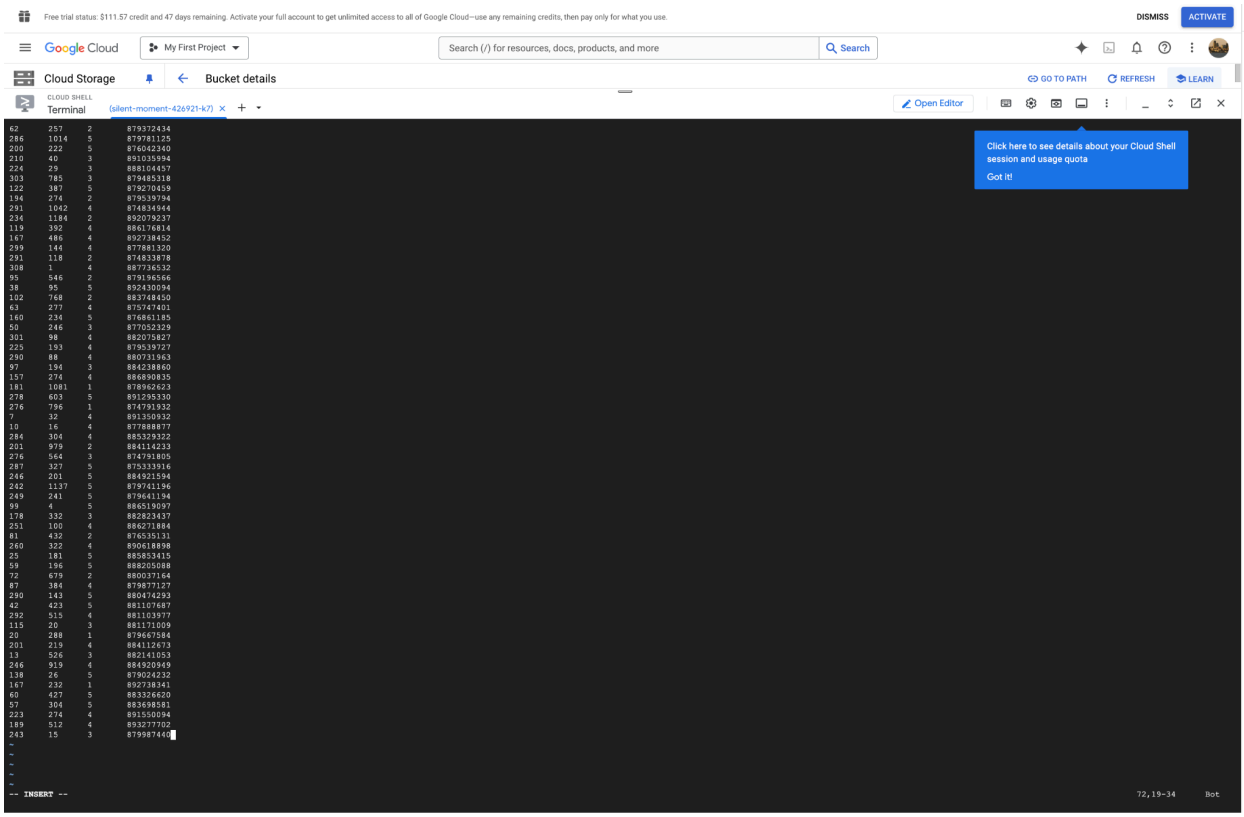
Free trial status: $111.57 credit and 47 days remaining. Activate your full account to get unlimited access to all of Google Cloud—use any remaining credits, then pay only for what you use.
DISMISS
ACTIVATE
Google Cloud
My First Project
Search (/) for resources, docs, products, and more
Search
Cloud Storage
Bucket details
GO TO PATH
REFRESH
LEARN
CLOUD SHELL
Terminal
(silent-moment-426921-k7)
Open Editor
Click here to see details about your Cloud Shell session and usage quota
Got it!
62 257 2 879372434
286 1014 5 879781125
200 222 5 876042340
210 40 3 891035994
224 29 3 888104457
303 785 3 879485318
122 387 5 879270459
184 274 2 879539794
291 1042 4 874834944
234 1184 2 892079237
119 392 4 886176814
167 486 4 892738452
299 144 4 877881320
291 118 2 874833878
308 1 4 887736532
95 546 2 879196566
38 95 5 892430094
102 768 2 883748450
63 277 4 875747401
160 234 5 876861185
50 246 3 877052329
301 98 4 882075827
225 193 4 879539727
290 88 4 880731963
97 194 3 884238860
157 274 4 886890835
181 1081 1 878962623
278 603 5 891295330
276 796 1 874791932
7 32 4 891350932
10 16 4 877888877
284 304 4 885329322
201 979 2 884114233
276 564 3 874791805
287 327 5 875333916
246 201 5 884921594
242 1137 5 879741196
249 241 5 879641194
99 4 5 886519097
178 332 3 882823437
251 100 4 886271884
81 432 2 876535131
260 322 4 890618898
25 181 5 885853415
59 196 5 888205088
72 679 2 880037164
87 384 4 879877127
290 143 5 880474293
42 423 5 881107687
292 515 4 881103977
115 20 3 881171009
20 288 1 879667584
201 219 4 884112673
13 526 3 882141053
246 919 4 884920949
138 26 5 879024232
167 232 1 892738341
60 427 5 883326620
57 304 5 883698581
223 274 4 891550094
189 512 4 893277702
243 15 3 879987440
~
~
~
~
~
-- INSERT --
72,19-34
Bot

```
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to silent-moment-426921-k7.
Use “gcloud config set project [PROJECT_ID]” to change to a different project.
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ vim u.data
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ echo '#!/bin/bash
> cat u.data | tr -s ' ' | cut -d' ' -f1-3 | tr ' ' ',' > u_data_transformed.csv' > transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ chmod +x transform_data.sh
```

```
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ ./transform_data.sh
tr: missing operand
Try 'tr --help' for more information.
cut: the delimiter must be a single character
Try 'cut --help' for more information.
tr: missing operand after ‘,’
Two strings must be given when translating.
Try 'tr --help' for more information.
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ cat u.data | tr -s ' ' | cut -d' ' -f1-3 | tr ' ' ',' > u_data_transformed.csv' > transform_data.sh
> ^C
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ cat u.data | tr -s ' ' | cut -d' ' -f1-3 | tr ' ' ',' >
u_data_transformed.csv' > transform_data.sh
-bash: syntax error near unexpected token `newline'
> u_data_transformed.csv' > transform_data.sh
-bash: $'u_data_transformed.csv > transform_data.sh\nu_data_transformed.csv': command not found
```

```
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ echo '#!/bin/bash
> cat u.data | tr -s ' ' | cut -d' ' -f1-3 | tr ' ' ',' > u_data_transformed.csv' > transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ chmod +x transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ ./transform_data.sh
tr: missing operand
Try 'tr --help' for more information.
cut: the delimiter must be a single character
Try 'cut --help' for more information.
tr: missing operand after ‘,’
Two strings must be given when translating.
Try 'tr --help' for more information.
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ echo '#!/bin/bash
> cat u.data | tr -s " " | cut -d" " -f1-3 | tr " " "," > u_data_transformed.csv' > transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ chmod +x transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ ./transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ echo '#!/bin/bash
> cat u.data | while read userid movieid rating timestamp
do
    echo "${userid},${movieid},${rating}"
done > u_data_transformed.csv' > transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ chmod +x transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ ./transform_data.sh
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gsutil cp u_data_transformed.csv gs://big_data_ml_recommendation_system/
Copying file://u_data_transformed.csv [Content-Type=text/csv]...
NotFoundException: 404 The destination bucket gs://big_data_ml_recommendation_system does not exist or the write to the destination must be restarted
```

```
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gsutil cp u_data_transformed.csv gs://big_data_recommendation_system/
Copying file://u_data_transformed.csv [Content-Type=text/csv]...
/ [1 files][  686.0 B/  686.0 B]
Operation completed over 1 objects/686.0 B.
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ vim recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ vim recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ rm recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ vim recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ rm recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ vim recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ rm recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ vim recommendation_example.py
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gsutil cp recommendation_example.py gs://big_data_ml_recommendation_system/
Copying file://recommendation_example.py [Content-Type=text/x-python]...
NotFoundException: 404 The destination bucket gs://big_data_ml_recommendation_system does not exist or the write to the destination must be restarted
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gsutil cp recommendation_example.py gs://big_data_recommendation_system/
Copying file://recommendation_example.py [Content-Type=text/x-python]...
/ [1 files][  1.0 KiB/  1.0 KiB]
Operation completed over 1 objects/1.0 KiB.
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gcloud dataproc jobs submit pyspark
gs://big_data_recommendation_system/recommendation_example.py \
 --cluster spark \
 --region us-central1
ERROR: (gcloud.dataproc.jobs.submit.pyspark) Exactly one of (--cluster | --cluster-labels) must be specified.
Usage: gcloud dataproc jobs submit pyspark PY_FILE (--cluster=CLUSTER | --cluster-labels=[KEY=VALUE,...]) [optional flags] [-- JOB_ARGS ...]
  optional flags may be  --archives | --async | --bucket | --cluster |
                         --cluster-labels | --driver-log-levels |
                         --driver-required-memory-mb |
                         --driver-required-vcores | --files | --help | --jars |
                         --labels | --max-failures-per-hour |
                         --max-failures-total | --properties |
                         --properties-file | --py-files | --region

For detailed information on this command and its flags, run:
  gcloud dataproc jobs submit pyspark --help
-bash: gs://big_data_recommendation_system/recommendation_example.py: No such file or directory
```

```
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gsutil cp recommendation_example.py gs://big_data_recommendation_system/
Copying file://recommendation_example.py [Content-Type=text/x-python]...
/ [1 files][  1.0 KiB/  1.0 KiB]
Operation completed over 1 objects/1.0 KiB.
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gcloud dataproc jobs submit pyspark
gs://big_data_recommendation_system/recommendation_example.py \
 --cluster spark \
 --region us-central1
ERROR: (gcloud.dataproc.jobs.submit.pyspark) Exactly one of (--cluster | --cluster-labels) must be specified.
Usage: gcloud dataproc jobs submit pyspark PY_FILE (--cluster=CLUSTER | --cluster-labels=[KEY=VALUE,...]) [optional flags] [-- JOB_ARGS ...]
  optional flags may be  --archives | --async | --bucket | --cluster |
                         --cluster-labels | --driver-log-levels |
                         --driver-required-memory-mb |
                         --driver-required-vcores | --files | --help | --jars |
                         --labels | --max-failures-per-hour |
                         --max-failures-total | --properties |
                         --properties-file | --py-files | --region

For detailed information on this command and its flags, run:
  gcloud dataproc jobs submit pyspark --help
-bash: gs://big_data_recommendation_system/recommendation_example.py: No such file or directory
```

```
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gcloud dataproc clusters create spark-cluster \
 --region us-west1 \
 --zone us-west1-a \
 --single-node
Waiting on operation [projects/silent-moment-426921-k7/regions/us-west1/operations/183268fc-3a5a-33d3-b0a7-be5cfc24882b].
Waiting for cluster creation operation...
WARNING: No image specified. Using the default image version. It is recommended to select a specific image version in production, as the default image version may change at any time.
WARNING: Consider using Auto Zone rather than selecting a zone manually. See https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/auto-zone
WARNING: Failed to validate permissions required for default service account: '235208717499-compute@developer.gserviceaccount.com'. Cluster creation could still be successful if required permissions have been granted to the respective service accou
nts as mentioned in the document https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/service-accounts#dataproc_service_accounts_2. This could be due to Cloud Resource Manager API hasn't been enabled in your project '235208717499' b
efore or it is disabled. Enable it by visiting 'https://console.developers.google.com/apis/api/cloudresourcemanager.googleapis.com/overview?project=235208717499'.
WARNING: The firewall rules for specified network or subnetwork would allow ingress traffic from 0.0.0.0/0, which could be a security risk.
WARNING: Unable to validate the staging bucket lifecycle configuration of the bucket 'dataproc-staging-us-west1-235208717499-e4xjsy3a' due to an internal error, Please make sure that the provided bucket doesn't have any delete rules set.
Waiting for cluster creation operation...done.
Created [https://dataproc.googleapis.com/v1/projects/silent-moment-426921-k7/regions/us-west1/clusters/spark-cluster] Cluster placed in zone [us-west1-a].
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$
```

Sign in to the gcloud CLI

You are seeing this page because you ran the following command in the gcloud CLI from this or another machine. If this is not the case, close this tab.

```
gcloud auth login --no-launch-browser
```

Enter the following verification code in gcloud CLI on the machine you want to log into. This is a credential **similar to your password** and should not be shared with others.

session and usage quota

Got it!

```
kadirbek_sharau@cloudshell:~ (silent-moment-426921-k7)$ gcloud dataproc jobs submit pyspark \
    gs://big_data_recommendation_system/recommendation_example.py \
    --cluster=spark-cluster \
    --region=us-west1
Job [a11ac11cd79947ef8200ae59818996a0] submitted.
Waiting for job output...
24/08/02 13:43:14 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker
24/08/02 13:43:14 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster
24/08/02 13:43:14 INFO org.apache.spark.SparkEnv: Registering BlockManagerMasterHeartbeat
24/08/02 13:43:14 INFO org.apache.spark.SparkEnv: Registering OutputCommitCoordinator
24/08/02 13:43:15 INFO org.sparkproject.jetty.util.log: Logging initialized @4149ms to org.sparkproject.jetty.util.log.Slf4jLog
24/08/02 13:43:15 INFO org.sparkproject.jetty.server.Server: jetty-9.4.40.v20210413; built: 2021-04-13T20:42:42.668Z; git: b881a572662e1943a14ae12e7e1207989f218b74; jvm 1.8.0_412-b08
24/08/02 13:43:15 INFO org.sparkproject.jetty.server.Server: Started @4257ms
24/08/02 13:43:15 INFO org.sparkproject.jetty.server.AbstractConnector: Started ServerConnector@2c3a691d{HTTP/1.1, (http/1.1)}{0.0.0.0:35311}
24/08/02 13:43:15 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at spark-cluster-m/10.138.0.15:8032
24/08/02 13:43:16 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at spark-cluster-m/10.138.0.15:10200
24/08/02 13:43:17 INFO org.apache.hadoop.conf.Configuration: resource-types.xml not found
24/08/02 13:43:17 INFO org.apache.hadoop.yarn.util.resource.ResourceUtils: Unable to find 'resource-types.xml'.
24/08/02 13:43:19 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1722605724517_0001
24/08/02 13:43:20 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at spark-cluster-m/10.138.0.15:8030
24/08/02 13:43:22 INFO com.google.cloud.hadoop.repackaged.gcs.com.google.cloud.hadoop.gcsio.GoogleCloudStorageImpl: Ignoring exception of type GoogleJsonResponseException; verified object already exists with desired state.
Traceback (most recent call last):
  File "/tmp/a11ac11cd79947ef8200ae59818996a0/recommendation_example.py", line 12, in <module>
    model = ALS.train(ratings, rank, numIterations)
  File "/usr/lib/spark/python/lib/pyspark.zip/pyspark/mllib/recommendation.py", line 280, in train
  File "/usr/lib/spark/python/lib/pyspark.zip/pyspark/mllib/recommendation.py", line 234, in _prepare
  File "/usr/lib/spark/python/lib/pyspark.zip/pyspark/rdd.py", line 1586, in first
  File "/usr/lib/spark/python/lib/pyspark.zip/pyspark/rdd.py", line 1533, in take
  File "/usr/lib/spark/python/lib/pyspark.zip/pyspark/rdd.py", line 2935, in getNumPartitions
  File "/usr/lib/spark/python/lib/py4j-0.10.9-src.zip/py4j/java_gateway.py", line 1304, in __call__
  File "/usr/lib/spark/python/lib/py4j-0.10.9-src.zip/py4j/protocol.py", line 326, in get_return_value
py4j.protocol.Py4JJavaError: An error occurred while calling o57.partitions.
: org.apache.hadoop.mapred.InvalidInputException: Input path does not exist: gs://big_data_ml_recommendation_sys/u_data_transformed.csv
        at org.apache.hadoop.mapred.LocatedFileStatusFetcher.getFileStatuses(LocatedFileStatusFetcher.java:156)
        at org.apache.hadoop.mapred.FileInputFormat.listStatus(FileInputFormat.java:247)
        at org.apache.hadoop.mapred.FileInputFormat.getSplits(FileInputFormat.java:325)
        at org.apache.spark.rdd.HadoopRDD.getPartitions(HadoopRDD.scala:205)
        at org.apache.spark.rdd.RDD.$anonfun$partitions$2(RDD.scala:300)
        at scala.Option.getOrElse(Option.scala:189)
        at org.apache.spark.rdd.RDD.partitions(RDD.scala:296)
        at org.apache.spark.rdd.MapPartitionsRDD.getPartitions(MapPartitionsRDD.scala:49)
        at org.apache.spark.rdd.RDD.$anonfun$partitions$2(RDD.scala:300)
        at scala.Option.getOrElse(Option.scala:189)
        at org.apache.spark.rdd.RDD.partitions(RDD.scala:296)
        at org.apache.spark.api.java.JavaRDDLike.partitions(JavaRDDLike.scala:61)
        at org.apache.spark.api.java.JavaRDDLike.partitions$(JavaRDDLike.scala:61)
        at org.apache.spark.api.java.AbstractJavaRDDLike.partitions(JavaRDDLike.scala:45)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
        at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:498)
        at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
        at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
        at py4j.Gateway.invoke(Gateway.java:282)
        at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
        at py4j.commands.CallCommand.execute(CallCommand.java:79)
        at py4j.GatewayConnection.run(GatewayConnection.java:238)
        at java.lang.Thread.run(Thread.java:750)

24/08/02 13:43:24 WARN org.apache.hadoop.util.concurrent.ExecutorHelper: Thread (Thread[GetFileInfo #0,5,main]) interrupted:
java.lang.InterruptedException
        at com.google.common.util.concurrent.AbstractFuture.get(AbstractFuture.java:510)
        at com.google.common.util.concurrent.FluentFuture$TrustedFuture.get(FluentFuture.java:88)
        at org.apache.hadoop.util.concurrent.ExecutorHelper.logThrowableFromAfterExecute(ExecutorHelper.java:48)
        at org.apache.hadoop.util.concurrent.HadoopThreadPoolExecutor.afterExecute(HadoopThreadPoolExecutor.java:90)
```